

LAB NO: 02

Task 1: Simple Linear Regression

1. Load dataset:

Use the Diabetes dataset from `sklearn.datasets`.

Select one feature (`bmi`) to predict the target (disease progression).

2. Perform Exploratory Data Analysis (EDA):

- o Plot scatter plot of BMI vs. Disease Progression.

- o Check correlation.

3. Implement Simple Linear Regression using `sklearn.linear_model.LinearRegression`:

- o Split data into training and testing sets.

- o Fit the model and predict disease progression.

- o Plot the regression line on the scatter plot.

4. Evaluate the model using:

- o Mean Squared Error (MSE)

- o R^2 score

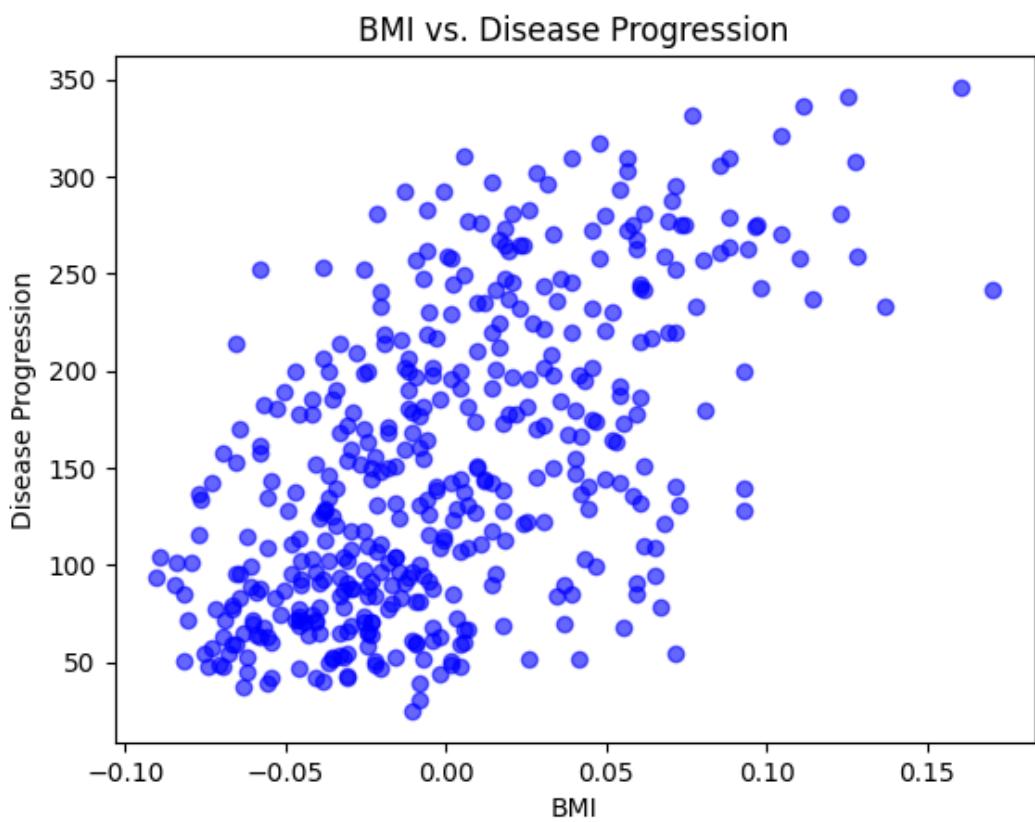
```
[ ]
```

```
▶ import pandas as pd
    import numpy as np
    from sklearn.datasets import load_diabetes
    from sklearn.linear_model import LinearRegression
    from sklearn.metrics import mean_squared_error, r2_score
    import seaborn as sns
```

```
diabetes = load_diabetes()
X = diabetes.data
y = diabetes.target
feature_names = diabetes.feature_names
```

```
X_bmi = X[:, np.newaxis, 2]
```

```
▶ plt.scatter(X_bmi, y, color="blue", alpha=0.6)
plt.xlabel("BMI")
plt.ylabel("Disease Progression")
plt.title("BMI vs. Disease Progression")
plt.show()
```



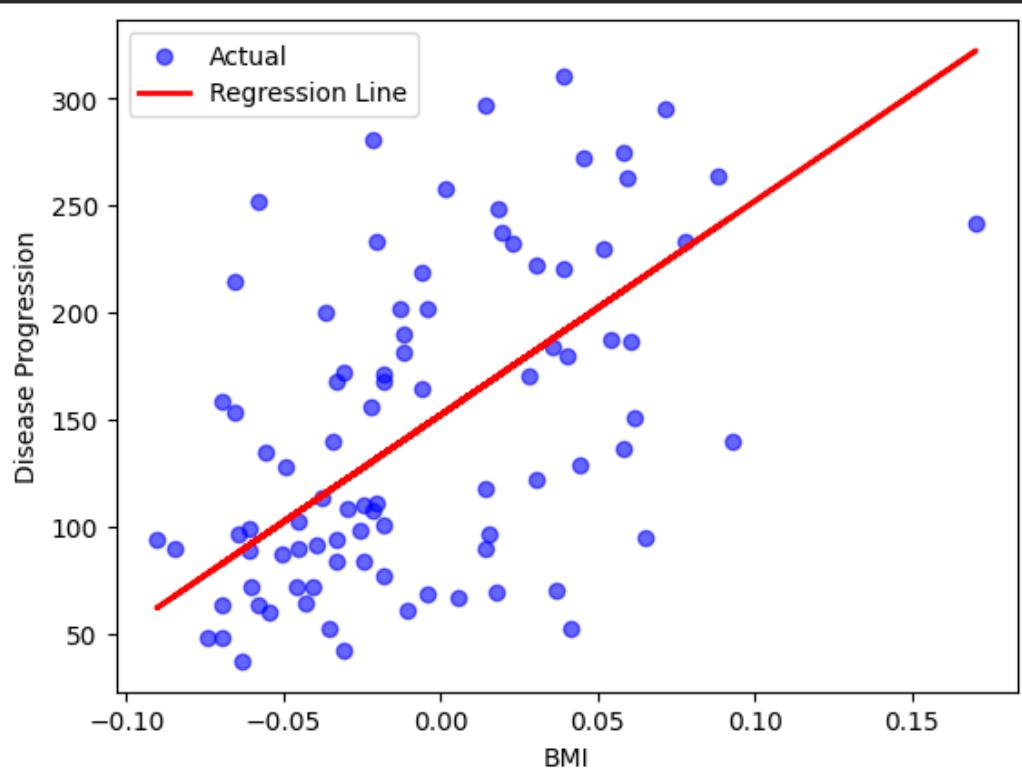
```
[ ] bmi_corr = np.corrcoef(X_bmi.flatten(), y)[0, 1]
print("Correlation between BMI and disease progression:", bmi_corr)

Correlation between BMI and disease progression: 0.5864501344746885

[ ] X_train, X_test, y_train, y_test = train_test_split(X_bmi, y, test_size=0.2, random_state=42)

[ ] model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

▶ plt.scatter(X_test, y_test, color="blue", alpha=0.6, label="Actual")
plt.plot(X_test, y_pred, color="red", linewidth=2, label="Regression Line")
plt.xlabel("BMI")
plt.ylabel("Disease Progression")
plt.legend()
plt.show()



```
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

```
print("MSE:", mse)  
print("R2 Score:", r2)
```

```
MSE: 4061.8259284949268  
R2 Score: 0.23335039815872138
```

▼ Task 2: Multivariate Linear Regression

1. Load dataset: Use the same Diabetes dataset, but include all 10 features to predict disease progression.

2. Perform EDA:

- o Generate a correlation heatmap between features and the target.
- o Create pair plots for selected features vs. target.

3. Implement Multivariate Linear Regression:

- o Use all independent variables to predict the target.
- o Fit and predict using the model.

4. Compare actual vs. predicted values using:

- o Scatter plot (predicted vs. actual).
- o Residual plot.

5. Evaluate model performance with:

- o MSE

- o RMSE

- o R² score

Which independent variable contributes the most to predicting disease progression?

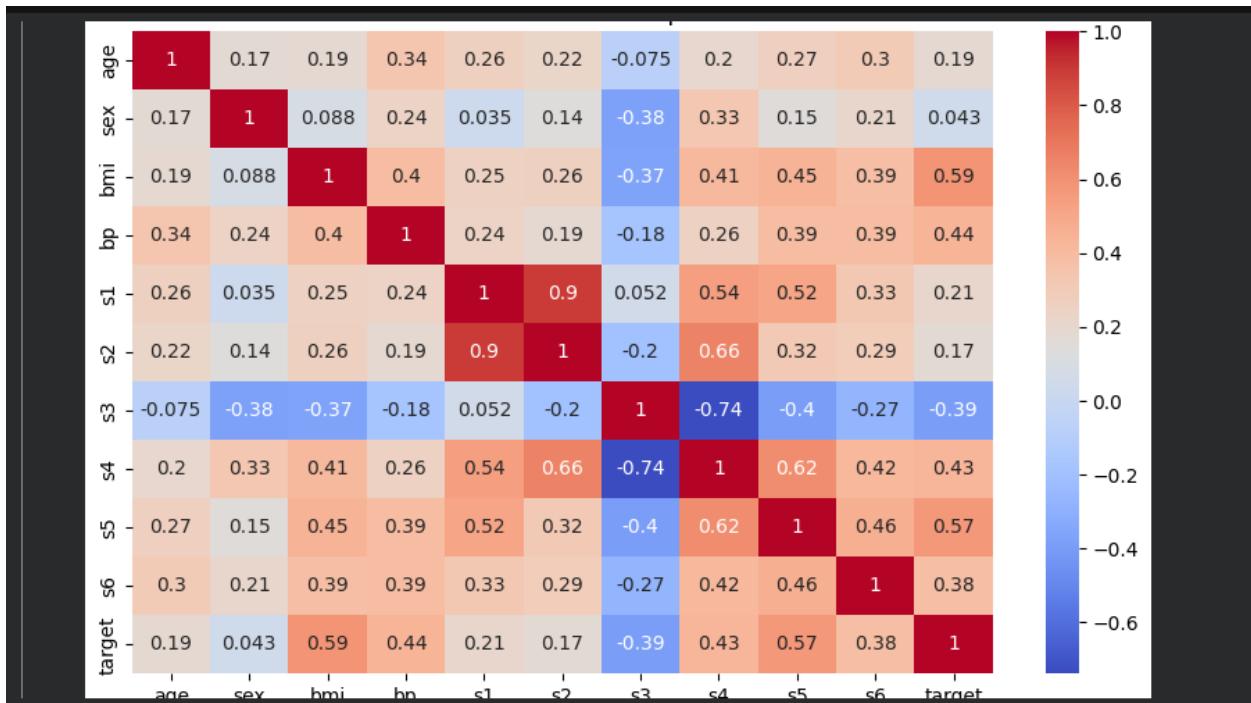
```

▶ # --- Use all features
X = diabetes.data
y = diabetes.target

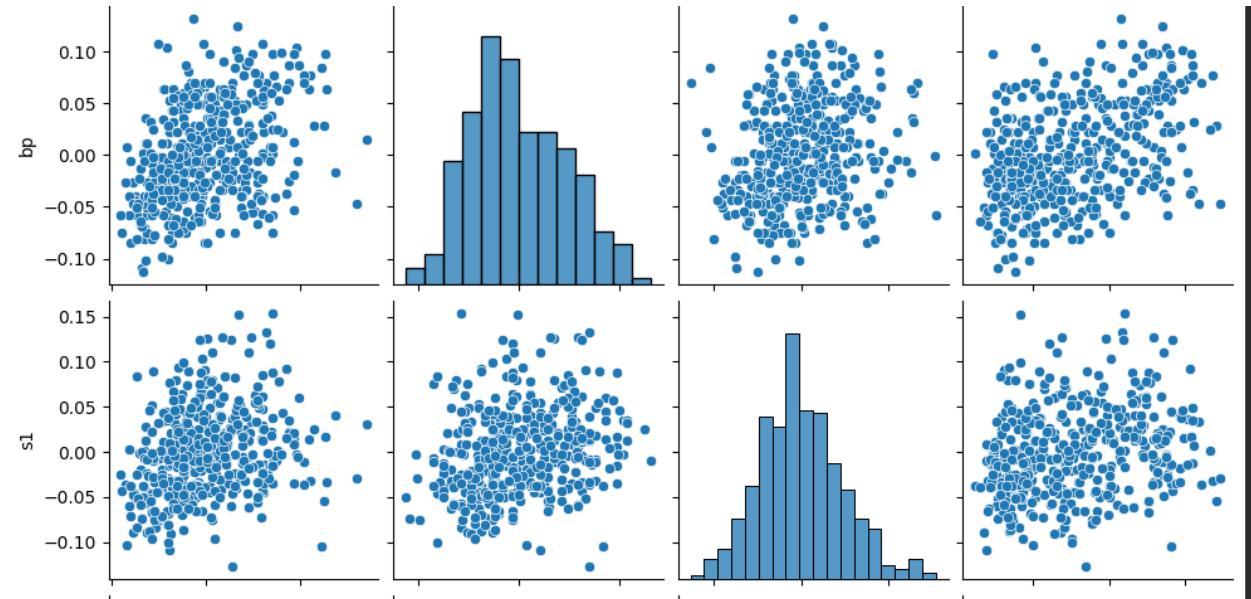
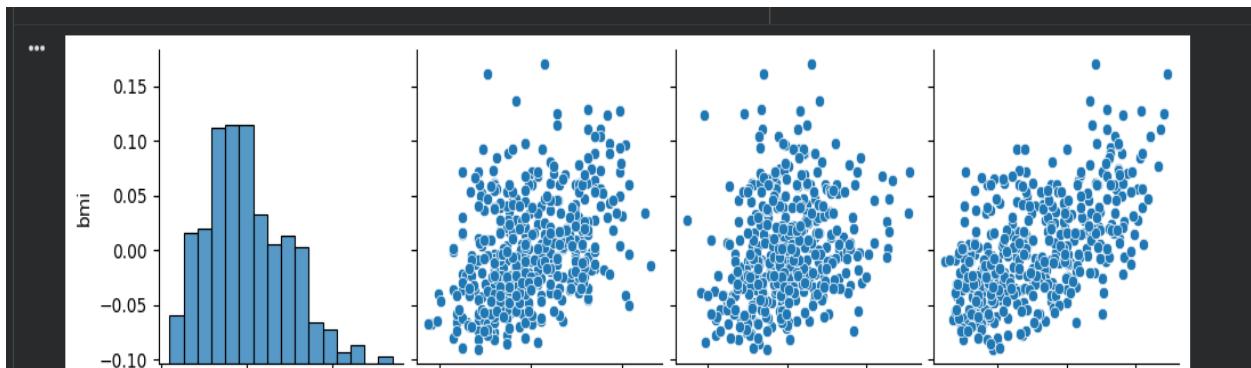
# --- EDA: Correlation Heatmap
df = pd.DataFrame(X, columns=feature_names)
df["target"] = y

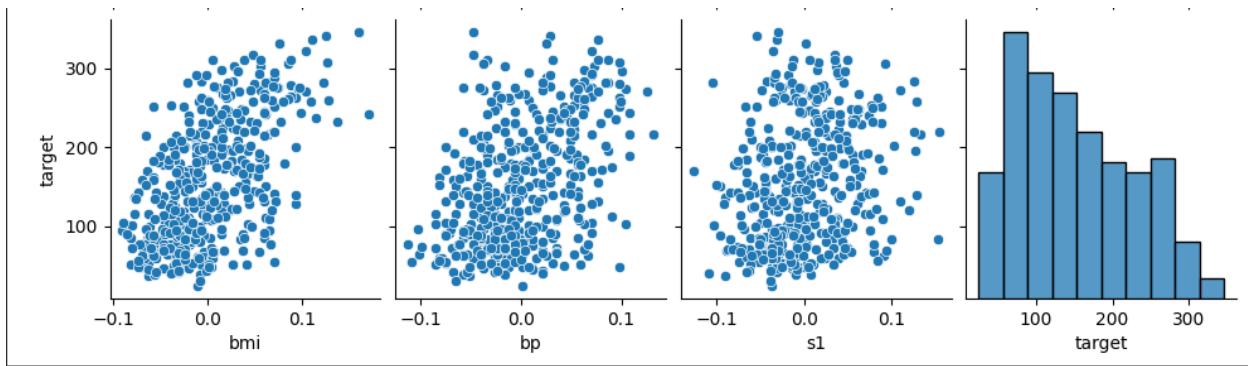
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()

```

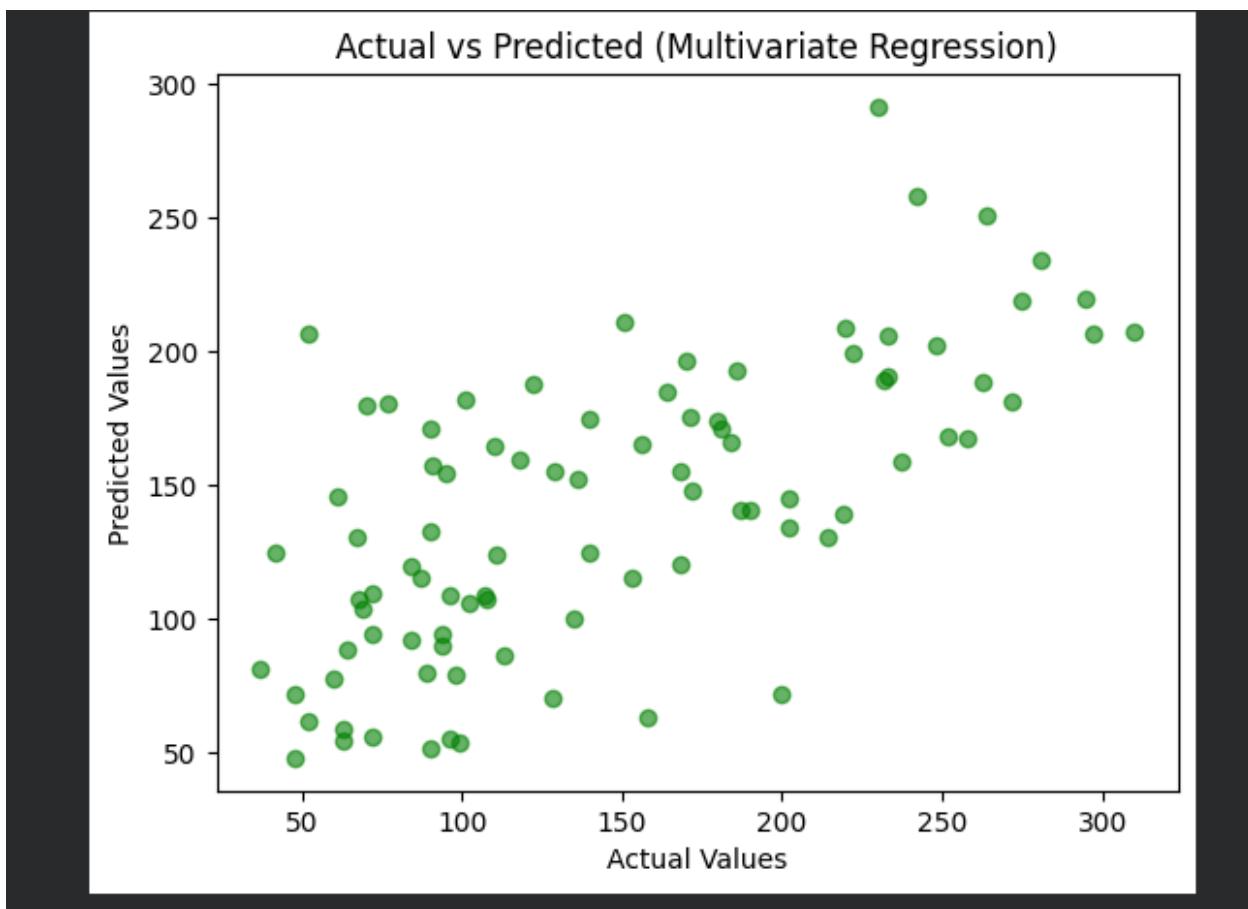


```
# --- Pairplot for selected features  
sns.pairplot(df[["bmi", "bp", "s1", "target"]])  
plt.show()  
  
# --- Train-Test Split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# --- Fit Multivariate Regression  
multi_model = LinearRegression()  
multi_model.fit(X_train, y_train)  
y_pred_multi = multi_model.predict(X_test)
```

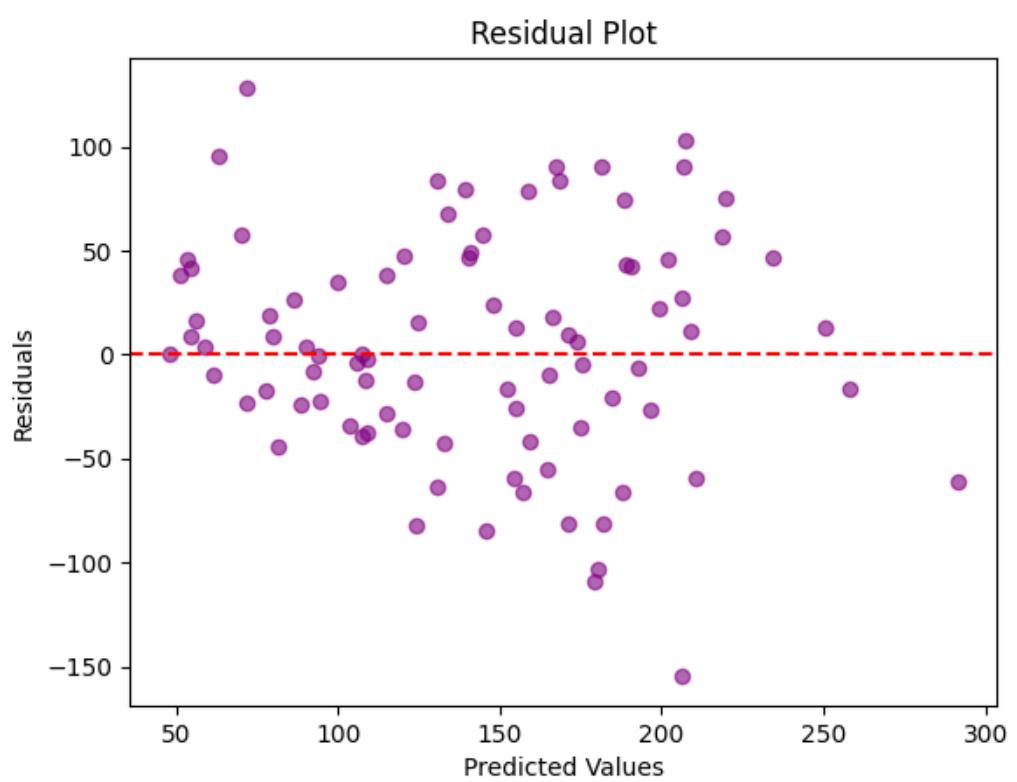




```
# --- Scatter: Actual vs Predicted
plt.scatter(y_test, y_pred_multi, color="green", alpha=0.6)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs Predicted (Multivariate Regression)")
plt.show()
```



```
# --- Residual Plot
residuals = y_test - y_pred_multi
plt.scatter(y_pred_multi, residuals, color="purple", alpha=0.6)
plt.axhline(y=0, color="red", linestyle="--")
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residual Plot")
plt.show()
```



```
# --- Evaluation
mse_multi = mean_squared_error(y_test, y_pred_multi)
rmse_multi = np.sqrt(mse_multi)
r2_multi = r2_score(y_test, y_pred_multi)

print("Multivariate MSE:", mse_multi)
print("Multivariate RMSE:", rmse_multi)
print("Multivariate R2 Score:", r2_multi)
```

```
Multivariate MSE: 2900.193628493482
Multivariate RMSE: 53.85344583676593
Multivariate R2 Score: 0.4526027629719195
```



```
# --- Find most contributing feature (by coefficient)
coefficients = pd.DataFrame({"Feature": feature_names, "Coefficient": multi_model.coef_})
coefficients = coefficients.sort_values(by="Coefficient", ascending=False)
print(coefficients)
```

	Feature	Coefficient
8	s5	736.198859
2	bmi	542.428759
5	s2	518.062277
3	bp	347.703844
7	s4	275.317902
6	s3	163.419983
9	s6	48.670657
0	age	37.904021
1	sex	-241.964362
4	s1	-931.488846

Task 3: Experimentation

1. Compare performance of simple vs. multivariate regression in terms of evaluation metrics.

[]

```
▶ print("\n--- Simple Linear Regression ---")
print("MSE:", mse)
print("R2 Score:", r2)

print("\n--- Multivariate Linear Regression ---")
print("MSE:", mse_multi)
print("RMSE:", rmse_multi)
print("R2 Score:", r2_multi)
```

▼

...

```
--- Simple Linear Regression ---
MSE: 4061.8259284949268
R2 Score: 0.23335039815872138
```

```
--- Multivariate Linear Regression ---
MSE: 2900.193628493482
RMSE: 53.85344583676593
R2 Score: 0.4526027629719195
```